


A • TARGET-HIDDEN INFERENCE

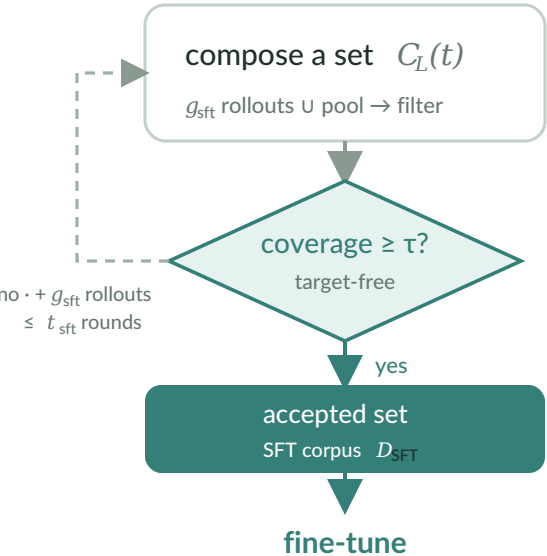
decompose within each rollout • compose across rollouts



the same verified clauses supervise both learning stages

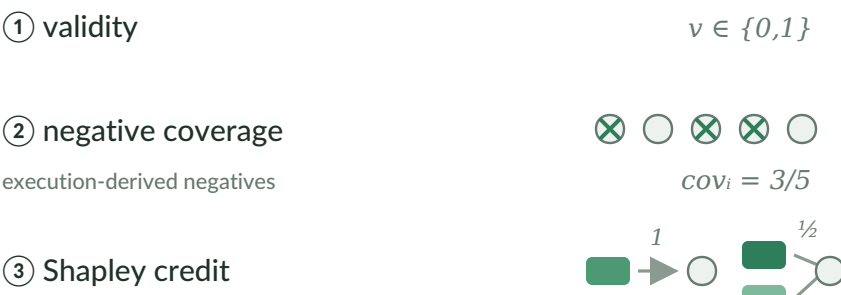
B • ALIGNED TRAINING

SFT • distill exploration



Target-free scorer

one target-free signal, two consumers



$$r_i = w_c \cdot cov_i + w_g \cdot \varphi_i - w_o \cdot o_i$$

RL • attribute exploration

